

Gradient descent updates parameters using the learning rate.
The update rule is $w = w - lr * \text{gradient}$.

Momentum accumulates past gradients to smooth the optimisation path.